

04 Numerical Optimization

May 29, 2021

Information: *Brief introduction to convexity, optimization, and gradient descent*

Written by: *Zihao Xu*

Last update date: *05.29.2021*

1 Convexity

1.1 Introduction

- Convexity plays a vital rule in the design of optimization algorithms, which is largely due to fact that it is much easier to analyze and test algorithms in such a context.
- If the algorithm performs poorly even in the convex setting, typically we should not hope to see great results otherwise.
- Even though the optimization problems in ML/DL are generally non-convex, they often exhibit some properties of convex ones near local minimum.

1.2 Open and Closed Sets

- Define

$$A \subset \mathbb{R}^n$$

and open ball of diameter ϵ is $B(r, \epsilon) = \{r \in \mathbb{R}^n : \|r - r_o\| < \epsilon\}$

- A set A is **open** if
 - At every point, there is an open ball contained in A
 - $\forall r \in A, \exists \epsilon > 0$ s.t. $B(r, \epsilon) \subset A$
- A set A is **closed** if $A^c = \mathbb{R}^n - A$ is open
- A set A is compact if it is closed and bounded
- Facts:
 - \mathbb{R}^N is both open and closed, but it is not compact
 - If A is compact, then every sequence in A has a limit point in A

1.3 Convex Sets

1.3.1 Definition

- A set C is convex if, for any $x, y \in C$ and $\theta \in \mathbb{R}$ with $0 \leq \theta \leq 1$:

$$\theta x + (1 - \theta)y \in C$$

- Intuitively, it means if we take any two elements in C and draw a line segment between these two elements, then every point on that line segment also belongs to C

- The point $\theta x + (1 - \theta)y$ is called a **convex combination** of the points x and y

1.3.2 Examples

- **All of \mathbb{R}^n .**

- Given any $\mathbf{x}, \mathbf{y} \in \mathbb{R}^n$,

$$\theta \mathbf{x} + (1 - \theta) \mathbf{y} \in \mathbb{R}^n$$

- **The non-negative orthant \mathbb{R}_+^n .**

- \mathbb{R}_+^n consists of all vectors in \mathbb{R}^n whose elements are all non-negative

$$\mathbb{R}_+^n = \{\mathbf{x} : x_i \geq 0 \ \forall i = 1, \dots, n\}$$

- Given any $\mathbf{x}, \mathbf{y} \in \mathbb{R}_+^N$ and $0 \leq \theta \leq 1$,

$$(\theta \mathbf{x} + (1 - \theta) \mathbf{y})_i = \theta x_i + (1 - \theta) y_i \geq 0 \ \forall i$$

- **Norm balls**

- Let $\|\cdot\|$ be some norm on \mathbb{R}^n (e.g., the Euclidean norm $\|\mathbf{x}\|_2 = \sqrt{\sum_{i=1}^n x_i^2}$). Then the set $\{\mathbf{x} : \|\mathbf{x}\| \leq 1\}$ is a convex set.

- Given $\mathbf{x}, \mathbf{y} \in \mathbb{R}^n$ with $\|\mathbf{x}\| \leq 1, \|\mathbf{y}\| \leq 1$ and $0 \leq \theta \leq 1$. Then

$$\|\theta \mathbf{x} + (1 - \theta) \mathbf{y}\| \leq \|\theta \mathbf{x}\| + \|(1 - \theta) \mathbf{y}\| = \theta \|\mathbf{x}\| + (1 - \theta) \|\mathbf{y}\| \leq 1$$

where the **triangle inequality** and the **positive homogeneity** of norms are used

- **Affine subspaces and polyhedra**

- Given a matrix $\mathbf{A} \in \mathbb{R}^{m \times n}$ and a vector $\mathbf{b} \in \mathbb{R}^m$, an affine subspace is the set $\{\mathbf{x} \in \mathbb{R}^n : \mathbf{Ax} = \mathbf{b}\}$ (note this could possibly be empty if \mathbf{b} is not in range of \mathbf{A}).

- Given $\mathbf{x}, \mathbf{y} \in \mathbb{R}^n$ s.t. $\mathbf{Ax} = \mathbf{Ay} = \mathbf{b}$, then for $0 \leq \theta \leq 1$:

$$\mathbf{A}(\theta \mathbf{x} + (1 - \theta) \mathbf{y}) = \theta \mathbf{Ax} + (1 - \theta) \mathbf{Ay} = \theta \mathbf{b} + (1 - \theta) \mathbf{b} = \mathbf{b}$$

- Similarly, a polyhedron is the set $\{\mathbf{x} \in \mathbb{R}^n : \mathbf{Ax} \preceq \mathbf{b}\}$ (also possibly empty), where \preceq denotes componentwise inequality

* All the entries of \mathbf{Ax} are less than or equal to their corresponding element in \mathbf{b}

- Given $\mathbf{x}, \mathbf{y} \in \mathbb{R}^n$ that satisfy $\mathbf{Ax} \leq \mathbf{b}$ and $\mathbf{Ay} \leq \mathbf{b}$ and $0 \leq \theta \leq 1$:

$$\mathbf{A}(\theta \mathbf{x} + (1 - \theta) \mathbf{y}) \leq \theta \mathbf{b} + (1 - \theta) \mathbf{b} = \mathbf{b}$$

- **Intersection of convex sets**

- Suppose C_1, C_2, \dots, C_k are convex sets. Then their intersection

$$\bigcap_{i=1}^k C_i = \{x : x \in C_i \ \forall i = 1, \dots, k\}$$

is also a convex set

- Given $x, y \in \bigcap_{i=1}^k C_i$ and $0 \leq \theta \leq 1$. Then

$$\theta x + (1 - \theta) y \in C_i \ \forall i = 1, \dots, k$$

by the definition of a convex set. Therefore

$$\theta x + (1 - \theta) y \in \bigcap_{i=1}^k C_i$$

- Note that the *union* of convex sets in general will not be convex
- **Positive semidefinite matrices**
 - The set of all symmetric positive semidefinite matrices, often times called the *positive semidefinite cone* and denoted \mathbb{S}_+^n is a convex set (in general, $\mathbb{S}^n \subset \mathbb{R}^{n \times n}$ denotes the set of symmetric $n \times n$ matrices).
 - A matrix $\mathbf{A} \in \mathbb{R}^{n \times n}$ is symmetric positive semidefinite if and only if $\mathbf{A} = \mathbf{A}^T$ and for all $\mathbf{x} \in \mathbb{R}^n$, $\mathbf{x}^T \mathbf{A} \mathbf{x} \geq 0$
 - Given two symmetric positive semidefinite matrices $\mathbf{A}, \mathbf{B} \in \mathbb{S}_+^n$ and $0 \leq \theta \leq 1$, then for any $\mathbf{x} \in \mathbb{R}^n$,
$$\mathbf{x}^T(\theta \mathbf{A} + (1 - \theta) \mathbf{B}) \mathbf{x} = \theta \mathbf{x}^T \mathbf{A} \mathbf{x} + (1 - \theta) \mathbf{x}^T \mathbf{B} \mathbf{x} \geq 0$$
 - The logic to show that all **positive definite**, **negative definite**, and **negative semidefinite** matrices are each also convex

1.4 Convex Functions

1.4.1 Definition

- A function $f : \mathbb{R}^n \rightarrow \mathbb{R}$ is convex if its domain (denoted $\mathcal{D}(f)$) is a *convex set*, and if, for all $\mathbf{x}, \mathbf{y} \in \mathcal{D}(f)$ and $\theta \in \mathbb{R}, 0 \leq \theta \leq 1$:

$$f(\theta \mathbf{x} + (1 - \theta) \mathbf{y}) \leq \theta f(\mathbf{x}) + (1 - \theta) f(\mathbf{y})$$

- Intuitively, it means if we pick any two points on the graph of a convex function and draw a straight line between them, then the portion of function between these two points will lie below this straight line.
- A function is called **strictly convex** if the definition holds with strict inequality for $\mathbf{x} \neq \mathbf{y}$ and $0 < \theta < 1$
- A function f is called **concave** if $-f$ is convex
- A function f is called **strictly concave** if $-f$ is strictly convex

1.4.2 First Order Condition for Convexity

- Suppose a function $f : \mathbb{R}^n \rightarrow \mathbb{R}$ is differentiable. Then f is convex if and only if $\mathcal{D}(f)$ is a convex set and for $\mathbf{x}, \mathbf{y} \in \mathcal{D}(f)$,

$$f(\mathbf{y}) \geq f(\mathbf{x}) + \nabla_x f(\mathbf{x})^T (\mathbf{y} - \mathbf{x})$$

where the function $f(\mathbf{x}) + \nabla_x f(\mathbf{x})^T (\mathbf{y} - \mathbf{x})$ is called the **first-order approximation** to the function f at the point \mathbf{x}

- Intuitively, this can be thought of as approximating f with its tangent line at the point \mathbf{x} .
- Similarly, f would be
 - strictly convex if this holds with strict inequality
 - concave if the inequality is reversed
 - strictly concave if the reverse inequality is strict

1.4.3 Second Order Condition for Convexity

- Suppose a function $f : \mathbb{R}^n \rightarrow \mathbb{R}$ is twice differentiable. Then f is convex if and only if $\mathcal{D}(f)$ is a convex set and its *Hessian* is positive semidefinite:

$$\forall \mathbf{x} \in \mathcal{D}(f), \nabla_{\mathbf{x}}^2 f(\mathbf{x}) \succeq 0$$

- Here the notation \succeq refers to positive semidefiniteness
- In one dimension, this is equivalent to the condition that the second derivative $f''(x)$ always be positive
- Similarly, f is
 - strictly convex if its Hessian is positive definite
 - concave if the Hessian is negative semidefinite
 - strictly concave if the Hessian is negative definite

1.4.4 Jensen's Inequality

- Start with the inequality in the basic definition of a convex function

$$f(\theta x + (1 - \theta)y) \leq \theta f(x) + (1 - \theta)f(y) \text{ for } 0 \leq \theta \leq 1$$

Using induction, extend this definition to convex combinations of more than one point

$$f\left(\sum_{i=1}^k \theta_i x_i\right) \leq \sum_{i=1}^k \theta_i f(x_i) \text{ for } \sum_{i=1}^k \theta_i = 1, \theta_i \geq 0 \forall i$$

This can also extend to infinite sums or integrals. In the latter case, the inequality can be written as

$$f\left(\int p(x) x dx\right) \leq \int p(x) f(x) dx \text{ for } \int p(x) dx = 1, p(x) \geq 0 \forall x$$

Since $\int p(x) dx = 1$, it is common to consider it a probability density, in which case the previous equation can be written in terms of expectations

$$f(\mathbb{E}[x]) \leq \mathbb{E}[f(x)]$$

which is called **Jensen's inequality**

1.4.5 Examples

- **Exponential**
 - Let $f : \mathbb{R} \rightarrow \mathbb{R}, f(x) = e^{ax}$ for any $a \in \mathbb{R}$.
 - $f''(x) = a^2 e^{ax}$ is positive for all x
- **Negative logarithm**
 - Let $f : \mathbb{R} \rightarrow \mathbb{R}, f(x) = -\log x$ with domain $\mathcal{D}(f) = \mathbb{R}_{++} = \{x : x > 0\}$
 - $f''(x) = \frac{1}{x^2} > 0$ for all x
- **Affine functions**
 - Let $f : \mathbb{R}^n \rightarrow \mathbb{R}, f(\mathbf{x}) = \mathbf{b}^T \mathbf{x} + c$ for some $\mathbf{b} \in \mathbb{R}^n, c \in \mathbb{R}$
 - The Hessian $\nabla_{\mathbf{x}}^2 f(\mathbf{x}) = 0$ for all \mathbf{x}
 - Affine functions of this form are the **only** functions that are **both convex and concave**
- **Quadratic function**
 - Let $f : \mathbb{R}^n \rightarrow \mathbb{R}, f(\mathbf{x}) = \frac{1}{2} \mathbf{x}^T \mathbf{A} \mathbf{x} + \mathbf{b}^T \mathbf{x} + c$ for a symmetric matrix $\mathbf{A} \in \mathbb{S}^n, \mathbf{b} \in \mathbb{R}^n$ and $c \in \mathbb{R}$
 - The Hessian for this function is $\nabla_{\mathbf{x}}^2 f(\mathbf{x}) = \mathbf{A}$
 - The convexity or non-convexity of f is determined entirely by whether or not \mathbf{A} is positive semidefinite

- The **squared Euclidean norm** $f(\mathbf{x}) = \|\mathbf{x}\|_2^2 = \mathbf{x}^T \mathbf{x}$ is a special case of quadratic functions where $\mathbf{A} = \mathbf{I}, \mathbf{b} = \mathbf{0}, c = 0$, so it is therefore a **strictly convex function**

- **Norms**

- Let $f : \mathbb{R}^n \rightarrow \mathbb{R}$ be some norm on \mathbb{R}^n .
- By the **triangle inequality** and **positive homogeneity** of norms, for $\mathbf{x}, \mathbf{y} \in \mathbb{R}^n, 0 \leq \theta \leq 1$,

$$f(\theta \mathbf{x} + (1 - \theta) \mathbf{y}) \leq f(\theta \mathbf{x}) + f((1 - \theta) \mathbf{y}) = \theta f(\mathbf{x}) + (1 - \theta) f(\mathbf{y})$$

- Not possible to prove convexity based on the first or second order conditions because norms are not generally differentiable

- **Nonnegative weighted sums of convex functions**

- Let f_1, f_2, \dots, f_k be convex functions and w_1, w_2, \dots, w_k be nonnegative real numbers. Then

$$f(x) = \sum_{i=1}^k w_i f_i(x)$$

is a convex function, since

$$\begin{aligned} f(\theta x + (1 - \theta)y) &= \sum_{i=1}^k w_i f_i(\theta x + (1 - \theta)y) \\ &\leq \sum_{i=1}^k w_i (\theta f_i(x) + (1 - \theta) f_i(y)) \\ &= \theta \sum_{i=1}^k w_i f_i(x) + (1 - \theta) \sum_{i=1}^k w_i f_i(y) \\ &= \theta f(x) + (1 - \theta) f(y) \end{aligned}$$

2 Optimization

2.1 Motivation

- Most ML/DL algorithms involve **optimization** of some sort.
 - Optimization refers to the task of either **minimizing** or maximizing some function $f(\mathbf{x})$ by altering \mathbf{x}
 - Usually phrase most optimization problems in terms of minimizing $f(\mathbf{x})$
 - Maximization may be accomplished via a minimization algorithm by minimizing $-f(\mathbf{x})$
- Usually the function we want to minimize is called the **objective function**, or **criterion**. In ML/DL contexts, the name **loss function** is often used.
 - As mentioned in introduction, a loss function quantifies the *distance* between the **real** and **predicted** value of the target.
 - Usually be a non-negative number where smaller values are better and perfect predictions incur a loss of 0
 - Usually denoted as $L(\theta)$ where θ is usually the parameter of ML/DL models
- Usually denote the value that minimizes a function with a superscript $*$
 - $\theta^* = \underset{\theta}{\operatorname{argmin}} L(\theta)$
- Most ML/DL algorithms are so complex that it is difficult or impossible to find the closed form solution for the optimization problem
 - Use numerical optimization method instead
- One common algorithm is **gradient descent**, other optimization algorithms are
 - Expectation Maximization
 - Sampling-based optimization
 - Greedy optimization

2.2 Local Minimum and Global Minimum

2.2.1 Local Minimum

- Let $f : A \rightarrow \mathbb{R}$ where $A \subset \mathbb{R}^N$, a point x is locally minimal if it is available and if there exists some $R > 0$ such that all feasible points z with $\|x - z\|_2 \leq R$, satisfy $f(x) \leq f(z)$
- **Necessary** condition for local minimum
 - Let f be continuously differentiable and let $x \in A$ be a local minimum, then $\nabla f(x) = 0$
- **Saddle Point**:
 - We say that $x \in A$ is a saddle point of f if $\nabla f(x) = 0$ and x is not a local minimum

2.2.2 Global Minimum

- A point x is globally minimal if it is available and for all feasible points z , $f(x) \leq f(z)$
 - A global minimum must also be a local minimum

2.3 Optimization Theorems

- Let $f : A \rightarrow \mathbb{R}$ where $A \subset \mathbb{R}^N$
 - If f is continuous and A is **compact**, then f takes on a global minimum in A
 - If f is **convex** on A , then any local minimum is a global minimum
 - If f is continuously differentiable and convex on A , then $\nabla f(x) = 0$ implies the $x \in A$ is a global minimum of f

- **Important Facts:**
 - Global minimum **may not be unique**
 - If A is closed but not bounded, then f may not take on a global minimum
 - Most interesting functions in ML/DL are **not** convex

2.4 Convex Optimization

- Formally, a convex optimization problem is an optimization problem of the form

$$\begin{array}{ll} \text{minimize} & f(x) \\ \text{subject to} & x \in C \end{array}$$

where f is a convex function, C is a convex set, and x is the optimization variable

- Often written as

$$\begin{array}{ll} \text{minimize} & f(x) \\ \text{subject to} & g_i(x) \leq 0 \quad i = 1, \dots, m \\ & h_i(x) = 0 \quad i = 1, \dots, p \end{array}$$

where f is a convex function, g_i are convex functions and h_i are affine functions and x is the optimization variable

2.5 Constrained Optimization

2.5.1 Definition

- Sometimes we wish not only to maximize or minimize a function $f(\mathbf{x})$ over all possible values of \mathbf{x} . Instead the maximal or minimal value of $f(\mathbf{x})$ for values of \mathbf{x} in some set \mathbb{S} . This is known as **constrained optimization**
- Points \mathbf{x} that lies within the set \mathbb{S} are called **feasible** points in constrained optimization terminology

2.5.2 Solution

- **Intuition:** Design a different, **unconstrained** optimization problem whose solution can be converted into a solution to the original constrained optimization problem
 - For example, to minimize $f(\mathbf{x})$ for $\mathbf{x} \in \mathbb{R}^2$ with \mathbf{x} constrained to have exactly unit L^2 norm, we can instead minimize

$$g(\theta) = f([\cos\theta, \sin\theta]^T)$$

with respect to θ , then return $[\cos\theta, \sin\theta]$ as the solution to the original problem

- Requires creativity
- The transformation between optimization problems must be designed specifically for each case we counter
- **Karush-Kuhn-Tucker (KKT)** approach
 - Provides a very general solution to constrained optimization

3 Gradient Descent

3.1 Basic Concepts

3.1.1 Definition

- **Definition:**
 - A **first-order iterative** optimization algorithm for finding **local minimum** of a **differential** function.
 - * The idea is to take *repeated steps* in the opposite direction of the *gradient* of the function at the current point, because this is the direction of steepest descent.
 - * As it only calculates the *first-order* derivative, it requires the objective function to be *differential* and is called *first-order optimization algorithms*
 - Some optimization algorithms that also use the Hessian matrix are called *second-order optimization algorithms*
 - * Converge when first-order derivative is zero, which only ensures reaching **local minimum** for general functions
 - That is to say, the start point will sometimes affect final convergence
 - * Generally speaking, gradient descent algorithms converge to the **global minimum** of continuously differentiable **convex** functions
 - **Theory:**
 - Based on the observation that if the multi-variable function $F(\mathbf{x})$ is defined and differentiable in a neighborhood of a point \mathbf{a} , then $F(\mathbf{x})$ decreases **fastest** if one goes from \mathbf{a} in the direction of the negative gradient of F at \mathbf{a} , which is $-\nabla F(\mathbf{a})$. It follows that if

$$\mathbf{a}_{n+1} = \mathbf{a}_n - \gamma \nabla F(\mathbf{a}_n)$$

for a $\gamma \in \mathbb{R}_+$ small enough, then

$$F(\mathbf{a}_n) \geq F(\mathbf{a}_{n+1})$$

- Simple form of **vanilla gradient descent** (GD):
 1. Start at random parameter $\boldsymbol{\theta}$
 2. Repeat until converged
 - $\mathbf{d} \leftarrow -\nabla L(\boldsymbol{\theta})$
 - $\boldsymbol{\theta} \leftarrow \boldsymbol{\theta} + \alpha \mathbf{d}^T$
 - α is called **learning rate** or **step size**

3.1.2 Compute Loss Gradient

- Take the **mean square error** as an example:

$$\begin{aligned}\nabla_{\boldsymbol{\theta}} L_{MSE}(\boldsymbol{\theta}) &= \nabla_{\boldsymbol{\theta}} \left\{ \frac{1}{N} \sum_{i=1}^N \|\mathbf{y}_i - f_{\boldsymbol{\theta}}(\mathbf{x}_i)\|^2 \right\} \\ &= \frac{1}{N} \sum_{i=1}^N \nabla_{\boldsymbol{\theta}} \{ (\mathbf{y}_i - f_{\boldsymbol{\theta}}(\mathbf{x}_i))^T (\mathbf{y}_i - f_{\boldsymbol{\theta}}(\mathbf{x}_i)) \}\end{aligned}$$

Use the chain rule and scale-by-vector matrix calculus identity that

$$\frac{\partial \mathbf{x}^T \mathbf{x}}{\partial \mathbf{x}} = 2\mathbf{x}^T$$

We can get

$$\begin{aligned}\nabla_{\boldsymbol{\theta}} L_{MSE}(\boldsymbol{\theta}) &= \frac{2}{N} \sum_{i=1}^N (\mathbf{y}_i - f_{\boldsymbol{\theta}}(\mathbf{x}_i))^T \nabla_{\boldsymbol{\theta}} (\mathbf{y}_i - f_{\boldsymbol{\theta}}(\mathbf{x}_i)) \\ &= \frac{2}{N} \sum_{i=1}^N (\mathbf{y}_i - f_{\boldsymbol{\theta}}(\mathbf{x}_i))^T \nabla_{\boldsymbol{\theta}} (-f_{\boldsymbol{\theta}}(\mathbf{x}_i)) \\ &= -\frac{2}{N} \sum_{i=1}^N (\mathbf{y}_i - f_{\boldsymbol{\theta}}(\mathbf{x}_i))^T \nabla_{\boldsymbol{\theta}} (f_{\boldsymbol{\theta}}(\mathbf{x}_i))\end{aligned}$$

- The result of the gradient usually includes three parts:
 - Sum over training data. It consists of a lot of computations but the way of computation is relatively easy and straight forward
 - Prediction error term such as $\mathbf{y}_i - f_{\boldsymbol{\theta}}(\mathbf{x}_i)$ in MSE, which is usually easy to get
 - Gradient of inference function $\nabla_{\boldsymbol{\theta}}(f_{\boldsymbol{\theta}}(\mathbf{x}_i))$, which is difficult to solve
 - * Enabled by **automatic differentiation** built into modern domain specific languages such as Pytorch, Tensorflow, ...
 - * For neural networks, this is known as **back propagation**

3.1.3 Select appropriate learning rate

- Too large α leads to instability and even divergence
- Too small α leads to slow convergence
- **Steepest gradient descent** use **line search** to compute the best α
 1. Start at random parameter $\boldsymbol{\theta}$
 2. Repeat until converged
 - $\mathbf{d} \leftarrow -\nabla L(\boldsymbol{\theta})$
 - $\alpha^* \leftarrow \underset{\alpha}{\operatorname{argmin}} \{L(\boldsymbol{\theta} + \alpha \mathbf{d}^T)\}$
 - $\boldsymbol{\theta} \leftarrow \boldsymbol{\theta} + \alpha^* \mathbf{d}^T$
- **Adaptive learning rates** may help, but not always
 - $\alpha = \frac{1}{t}$, approaches 0 but can cover an infinite distance since $\lim_{a \rightarrow \infty} \sum_{t=1}^a \frac{1}{t} = \infty$
- **Coordinate Descent** update one parameter at a time
 - Removes problem of selecting step size
 - Each update can be very fast, but lots of updates

3.1.4 Slow convergence due to Poor Conditioning

- **Conditioning** refers to how rapidly a function changes with respect to small changes in its inputs.
- Consider the function

$$f(x) = \mathbf{A}^{-1} \mathbf{x}$$

When $\mathbf{A} \in \mathbb{R}^{n \times n}$ has an eigenvalue decomposition, its **condition number** is

$$\max_{i,j} \left| \frac{\lambda_i}{\lambda_j} \right|$$

This is the ratio of the magnitude of the largest and smallest eigenvalue

- A problem with a **low condition number** is said to be **well-conditioned**, while a problem with a high condition number is said to be ill-conditioned

- In non-mathematical terms, an ill-conditioned problem is one where, for a small change in the inputs there is a large change in the answer or dependent variable, which means the correct solution to the equation becomes hard to find
- Condition number is a property of the problem
- **Gradient descent** is very sensitive to **condition number** of the problem
 - No good choice of step size. Tiny change in one variable could lead to great change in dependent variable.
- **Solutions:**
 - **Newton's method:** Correct for local second derivative.
 - * Too much computation and too difficult to implement
 - * Harmful when near saddle points
 - **Alternative methods:**
 - * Preconditioning: Easy, but tends to be ad-hoc, not so robust
 - * Momentum

3.1.5 Vanishing Gradients

- The most insidious problem to encounter
- Some function leads to almost zero gradients far away from local minimums, which makes the optimization stuck for a long time or even stop.
- For example, assume that we want to minimize the function

$$f(x) = \tanh(x)$$

The derivative is

$$f'(x) = 1 - \tanh^2(x)$$

If we happen to get started at $x = 4$ then the derivative at that point is

$$f'(4) = 0.0013$$

The gradient is close to nil. Consequently, optimization will get stuck for a long time before we make progress

- **Possible Solutions:**
 - Reparameterize the problem
 - Good initialization of the parameter
 - Reconstruct the objective function (e.g., change activation function in neural networks)