

05 AutoGrad in PyTorch

June 9, 2021

Information: *Brief introduction to Tensor in PyTorch, automatic differentiation theory, how to use AutoGrad in PyTorch and two simple examples*

Written by: *Zihao Xu*

Last update date: *06.09.2021*

1 Data Manipulation via PyTorch

- The n -dimensional array is usually called the tensor
- *tensor* class in Pytorch is similar to *NumPy*'s *ndarray* with several additional features
 - GPU is well-supported to accelerate the computation whereas *NumPy* only supports CPU computation
 - *tensor* class supports automatic differentiation
- The [Pytorch documentation](#) shows the full attributes

1.1 Create a tensor

- To get started, import **torch**. Although it's called PyTorch, we should import **torch** instead of **pytorch**

```
[1]: import torch
      # Check the version of a module
      print(torch.__version__)
```

1.8.1

- Common ways to creating a tensor
 - `torch.arange(start,end,step)`
 - `torch.zeros(shape)`
 - `torch.ones(shape)`
 - `torch.randn(shape)`
 - `torch.tensor(elements)`

```
[2]: torch.arange(0, 12, 1)
```

```
[2]: tensor([ 0,  1,  2,  3,  4,  5,  6,  7,  8,  9, 10, 11])
```

```
[3]: torch.zeros((3, 4))
```

```
[3]: tensor([[0., 0., 0., 0.],
            [0., 0., 0., 0.],
            [0., 0., 0., 0.]])
```

```
[4]: torch.ones((2, 5))
```

```
[4]: tensor([[1., 1., 1., 1., 1.],
            [1., 1., 1., 1., 1.]])
```

```
[5]: torch.randn(10)
```

```
[5]: tensor([ 0.1626, -0.9563,  1.0995,  0.8781,  0.5974,  0.1148,  0.2608,  1.3726,
            0.8922, -0.6782])
```

```
[6]: torch.tensor([[1, 2, 3], [4, 5, 6]])
```

```
[6]: tensor([[1, 2, 3],
            [4, 5, 6]])
```

- One can access a tensor's shape by viewing the **shape** attribute

```
[7]: torch.ones((4, 5)).shape
```

```
[7]: torch.Size([4, 5])
```

- **reshape** method can change the shape of a tensor without altering either the number of elements or their values.
 - No need to manually specify every dimension
 - Can place `-1` for the dimension that we would like tensors to automatically infer

```
[8]: torch.arange(12).reshape(3, 4)
```

```
[8]: tensor([[ 0,  1,  2,  3],
            [ 4,  5,  6,  7],
            [ 8,  9, 10, 11]])
```

```
[9]: torch.arange(12).reshape(3, -1)
```

```
[9]: tensor([[ 0,  1,  2,  3],
            [ 4,  5,  6,  7],
            [ 8,  9, 10, 11]])
```

1.2 Type of a tensor

- Usually, a tensor is created as **tensor.float32** (32-bit floating point) by default. One can view its type in the **dtype** attribute
 - When creating a tensor using **torch.tensor**, tensor with all integers would be created as **torch.int64** (64-bit signed integer)
- Full tensor types can be viewed in the [documentation](#)

```
[10]: torch.ones(10).dtype
```

```
[10]: torch.float32
```

```
[11]: torch.tensor([[1, 2, 3], [4, 5, 6]]).dtype
```

```
[11]: torch.int64
```

```
[12]: # Only add one dot after the first element  
torch.tensor([[1., 2, 3], [4, 5, 6]]).dtype
```

```
[12]: torch.float32
```

- One can assign the wanted type when creating the tensor by setting the **dtype** attribute to
 - *torch.float32*, 32-bit floating point
 - *torch.float64*, 64-bit floating point
 - *torch.uint8*, 8-bit unsigned integer
 - *torch.int8*, 8-bit signed integer
 - *torch.int32*, 32-bit signed integer
 - *torch.int64*, 64-bit signed integer
 - *torch.bool*, Boolean

```
[13]: torch.ones(10, dtype=torch.float64).dtype
```

```
[13]: torch.float64
```

```
[14]: torch.ones(10, dtype=torch.uint8).dtype
```

```
[14]: torch.uint8
```

```
[15]: torch.ones(10, dtype=torch.int32).dtype
```

```
[15]: torch.int32
```

- One can also construct the type of a tensor from list or numpy array using the method:
 - *FloatTensor*, 32-bit floating point
 - *DoubleTensor*, 64-bit floating point
 - *ByteTensor*, 8-bit unsigned integer
 - *CharTensor*, 8-bit signed integer
 - *IntTensor*, 32-bit signed integer
 - *LongTensor*, 64-bit signed integer
 - *BoolTensor*, Boolean

```
[16]: torch.DoubleTensor([1, 2, 3]).dtype
```

```
[16]: torch.float64
```

```
[17]: torch.ByteTensor([1, 2, 3]).dtype
```

```
[17]: torch.uint8
```

1.3 Use GPU for tensor computation

- Unless otherwise specified, a new tensor will be stored in main memory and designated for CPU-based computation
- One can check which device the tensor is designated for by viewing the **device** attribute

```
[18]: torch.ones(10).device
```

```
[18]: device(type='cpu')
```

- One can always set the create a device if a GPU supporting cuda is available and use **to(device)** method to determine the device on which a tensor is or will be allocated
 - Assign the **device** parameter when creating a tensor also works

```
[19]: cuda0 = torch.device("cuda:0" if torch.cuda.is_available() else "cpu")
      cuda0
```

```
[19]: device(type='cuda', index=0)
```

```
[20]: torch.ones(5, device=cuda0)
```

```
[20]: tensor([1., 1., 1., 1., 1.], device='cuda:0')
```

```
[21]: # If available, indexing the cuda also works
      torch.ones(5, device=0)
```

```
[21]: tensor([1., 1., 1., 1., 1.], device='cuda:0')
```

```
[22]: # If available, the string also works
      torch.ones(5, device="cuda:0")
```

```
[22]: tensor([1., 1., 1., 1., 1.], device='cuda:0')
```

```
[23]: # Use the to method the move a tensor
      x = torch.ones(5).to(cuda0)
      x
```

```
[23]: tensor([1., 1., 1., 1., 1.], device='cuda:0')
```

```
[24]: # One can also move a tensor from GPU to CPU
      x.to("cpu").device
```

```
[24]: device(type='cpu')
```

1.4 Operations

- Common standard arithmetic operators have all been lifted to element-wise operations

```
[25]: x = torch.tensor([1., 2., 4., 8.])
      c = 1.
      x + c, x * c, x**c
```

```
[25]: (tensor([2., 3., 5., 9.]), tensor([1., 2., 4., 8.]), tensor([1., 2., 4., 8.]))
```

```
[26]: x = torch.tensor([1., 2., 4., 8.])
      y = torch.tensor([4., 3., 2., 1.])
      x + y, x * y, x**y
```

```
[26]: (tensor([5., 5., 6., 9.]),
      tensor([4., 6., 8., 8.]),
      tensor([ 1.,  8., 16.,  8.]))
```

```
[27]: torch.exp(x)
```

```
[27]: tensor([2.7183e+00, 7.3891e+00, 5.4598e+01, 2.9810e+03])
```

- Matrix multiplication is also supported

```
[28]: A = torch.randn((4, 3))
      B = torch.randn((3, 5))
      # Two ways of matrix multiplication
      torch.mm(A, B), A @ B
```

```
[28]: (tensor([[ -1.9090,  4.4974, -0.7523,  2.8361, -2.4656],
              [-0.7529,  0.9086,  0.8353,  0.1293,  0.1171],
              [-0.7998,  1.4260,  2.3517,  1.2654,  0.0200],
              [ 2.2119, -5.1211, -1.6146, -3.8720,  2.1983]]),
      tensor([[ -1.9090,  4.4974, -0.7523,  2.8361, -2.4656],
              [-0.7529,  0.9086,  0.8353,  0.1293,  0.1171],
              [-0.7998,  1.4260,  2.3517,  1.2654,  0.0200],
              [ 2.2119, -5.1211, -1.6146, -3.8720,  2.1983]]))
```

```
[29]: A = torch.randn((5, 4))
      B = torch.randn((5, 4))
      # Two ways of elementwise multiplication
      torch.mul(A, B), A * B
```

```
[29]: (tensor([[ -1.5337,  0.5801, -0.0622,  0.1134],
              [-0.0442, -0.2856, -0.1522,  0.0301],
              [-1.4391, -0.2859, -0.1268,  0.1497],
              [ 0.4119, -1.3713, -0.1635,  0.3135],
              [ 0.3121, -0.7067, -0.0231, -2.0135]]),
      tensor([[ -1.5337,  0.5801, -0.0622,  0.1134],
              [-0.0442, -0.2856, -0.1522,  0.0301],
              [-1.4391, -0.2859, -0.1268,  0.1497],
              [ 0.4119, -1.3713, -0.1635,  0.3135],
              [ 0.3121, -0.7067, -0.0231, -2.0135]]))
```

```
[ 0.3121, -0.7067, -0.0231, -2.0135]]))
```

- We can also **concatenate** multiple tensors together, stacking them end-to-end to form a larger tensor. We just need to provide a list of tensors and tell the system along which axis to concatenate.

```
[30]: A = torch.arange(0, 12, 1).reshape((3, 4))
      B = torch.ones((3, 4))
      # dim stands for the index of dimension in which the tensors are concatenated
      torch.cat((A, B), dim=0), torch.cat((A, B), dim=1)
```

```
[30]: (tensor([[ 0.,  1.,  2.,  3.],
               [ 4.,  5.,  6.,  7.],
               [ 8.,  9., 10., 11.],
               [ 1.,  1.,  1.,  1.],
               [ 1.,  1.,  1.,  1.],
               [ 1.,  1.,  1.,  1.]]),
      tensor([[ 0.,  1.,  2.,  3.,  1.,  1.,  1.,  1.],
               [ 4.,  5.,  6.,  7.,  1.,  1.,  1.,  1.],
               [ 8.,  9., 10., 11.,  1.,  1.,  1.,  1.])))
```

- Also, we can construct a binary tensor via logical statements

```
[31]: A = torch.arange(0, 12, 1).reshape((3, 4))
      B = 5 * torch.ones((3, 4))
      A == B, A > B
```

```
[31]: (tensor([[False, False, False, False],
               [False, True, False, False],
               [False, False, False, False]]),
      tensor([[False, False, False, False],
               [False, False, True, True],
               [ True,  True,  True,  True]]))
```

1.5 Broadcasting Mechanism

- Under certain conditions, even shapes differ, we can still perform **element-wise** operations by invoking the **broadcasting mechanism**.
 - First, expand one or both arrays by copying elements appropriately so that after this transformation, the two tensors have the same shape.
 - Second, carry out the element-wise operation on the resulting arrays
- In most cases, we broadcast along an axis where an array initially only has length 1

```
[32]: a = torch.arange(3).reshape(3, 1)
      b = torch.arange(2).reshape(1, 2)
      a, b, a + b
```

```
[32]: (tensor([[0],
               [1],
```

```

        [2]]),
tensor([[0, 1]]),
tensor([[0, 1],
        [1, 2],
        [2, 3]]))

```

1.6 Indexing and Slicing

- As in standard Python lists, we can access elements according to their relative position to the end of the list by using negative indices

```

[33]: X = torch.arange(16).reshape(4, 4)
      X, X[-1], X[1:3]

```

```

[33]: (tensor([[ 0,  1,  2,  3],
               [ 4,  5,  6,  7],
               [ 8,  9, 10, 11],
               [12, 13, 14, 15]]),
      tensor([12, 13, 14, 15]),
      tensor([[ 4,  5,  6,  7],
               [ 8,  9, 10, 11]]))

```

- Can also index using binary tensor

```

[34]: A = torch.arange(0, 12, 1).reshape((3, 4))
      B = 5 * torch.ones((3, 4))
      A, B, A[A > B]

```

```

[34]: (tensor([[ 0,  1,  2,  3],
               [ 4,  5,  6,  7],
               [ 8,  9, 10, 11]]),
      tensor([[5., 5., 5., 5.],
               [5., 5., 5., 5.],
               [5., 5., 5., 5.]]),
      tensor([ 6,  7,  8,  9, 10, 11]))

```

2 Automatic Calculation of Gradients

- In practice, based on our designed model, the system builds a **computational graph**, tracking which data combined through which operations to produce the output. Automatic differentiation enables the system to subsequently **backpropagate gradients**.
 - Here *backpropagate* simply means to trace through the computational graph, filling in the partial derivatives with respect to each parameter

2.1 Intuition

- All computation can be broken into simple components
 - sum
 - multiply
 - exponential
 - convolution
 - ...
- Derivatives for each simple component can be derived mathematically
- Derivatives for **any composition** can be derived via **chain rule**

2.2 Automatic Differentiation - Forward Mode

- **Notice:** Here is only my naive understanding about automatic differentiation since I haven't found some detailed mathematical explanation about it.
- One naive way of find the partial derivative respect to i -th variable $\frac{\partial f}{\partial x_i}$ in the function

$$f = f(x_1, x_2, \dots, x_n)$$

the simplest way is to **construct** the chain rule ($\frac{\partial f}{\partial t}$ is actually meaningless in this context)

$$\frac{\partial f}{\partial t} = \sum_{j=1}^n \frac{\partial f}{\partial x_j} \frac{\partial x_j}{\partial t}$$

by setting

$$\frac{\partial x_j}{\partial t} = \begin{cases} 1 & j = i \\ 0 & j \neq i \end{cases}$$

we get

$$\frac{\partial f}{\partial t} = \frac{\partial f}{\partial x_i}$$

- That is to say, for each **decomposed simple operation**, the computer does not only calculate the **output**, but also tracks the **derivatives of output respect to time** based on the input value and input derivative, thus **step by step** getting the final derivative respect to time, which is equal to the partial derivative
- Since the computation is **straight forward**, this method to calculate partial derivative is called **forward mode**.
- For example, for a function

$$f(x_1, x_2) = \ln(x_1) + x_1 x_2 - \sin(x_2)$$

Decompose the function into simple operation nodes


$$\begin{aligned}v_{-1} &= x_1 \\ v_0 &= x_2 \\ v_1 &= \ln v_{-1} \\ v_2 &= v_{-1} \times v_0 \\ v_3 &= \sin v_0 \\ v_4 &= v_1 + v_2 \\ v_5 &= v_4 - v_3 \\ y &= v_5\end{aligned}$$

The forward mode computation graph is shown below

```
[35]: from graphviz import Digraph

f = Digraph('ComputationGraph')
f.attr(rankdir='LR')
f.attr('node', shape='circle')
f.node('v_{-1}', label='<v<sub>-1</sub>>')
f.node('v_0', label='<v<sub>0</sub>>')
f.node('v_1', label='<v<sub>1</sub>>')
f.node('v_2', label='<v<sub>2</sub>>')
f.node('v_3', label='<v<sub>3</sub>>')
f.node('v_4', label='<v<sub>4</sub>>')
f.node('v_5', label='<v<sub>5</sub>>')
f.edge('v_{-1}', 'v_1')
f.edge('v_{-1}', 'v_2')
f.edge('v_0', 'v_2')
f.edge('v_0', 'v_3')
f.edge('v_1', 'v_4')
f.edge('v_2', 'v_4')
f.edge('v_4', 'v_5')
f.edge('v_3', 'v_5')
f
```

[35]:



- To find the **partial derivate** $\frac{\partial f}{\partial x_1}$ when $x_1 = 2, x_2 = 5$, compute along the **forward evaluation trace** first:

$$\begin{array}{lll}
v_{-1} = & x_1 = & 2 \\
v_0 = & x_2 = & 5 \\
v_1 = & \ln v_{-1} = & \ln 2 \\
v_2 = & v_{-1} \times v_0 = & 2 \times 5 \\
v_3 = & \sin v_0 = & \sin 5 \\
v_4 = & v_1 + v_2 = & 0.693 + 10 \\
v_5 = & v_4 - v_3 = & 10.693 + 0.959 \\
y = & v_5 = & 11.652
\end{array}$$

Then set $\dot{x}_1 = 1, \dot{x}_2 = 0$ and compute along the **forward derivative trace** (maybe also computed simultaneously along with the forward evaluation trace)

$$\begin{array}{lll}
\dot{v}_{-1} = & \dot{x}_1 = & 1 \\
\dot{v}_0 = & \dot{x}_2 = & 0 \\
\dot{v}_1 = & \dot{v}_{-1}/v_{-1} = & 1/2 \\
\dot{v}_2 = & \dot{v}_{-1} \times v_0 + \dot{v}_0 \times v_{-1} = & 1 \times 5 + 0 \times 2 \\
\dot{v}_3 = & \dot{v}_0 \times \cos v_0 = & 0 \times \cos 5 \\
\dot{v}_4 = & \dot{v}_1 + \dot{v}_2 = & 0.5 + 5 \\
\dot{v}_5 = & \dot{v}_4 - \dot{v}_3 = & 5.5 - 0 \\
\dot{y} = & \dot{v}_5 = & 5.5
\end{array}$$

Therefore, the result is

$$\frac{\partial f}{\partial x_1} = \frac{\partial f}{\partial t} = 5.5$$

If $\frac{\partial f}{\partial x_2}$ is also interested, the **forward derivative trace** needs to be computed again when setting $\dot{x}_1 = 0, \dot{x}_2 = 1$

- **Comments:**
 - For a function $f : \mathbb{R}^n \Rightarrow \mathbb{R}^m$ with input $\mathbf{x} \in \mathbb{R}^n$ and output $\mathbf{y} \in \mathbb{R}^m$, needs to calculate n times along the **forward derivative trace** for the partial derivatives

$$\frac{\partial \mathbf{y}}{\partial x_j}, \quad j = 1, 2, \dots, n$$

by sequentially setting

$$\frac{\partial x_j}{\partial t} = \begin{cases} 1 & j = i \\ 0 & j \neq i \end{cases}, \quad i = 1, 2, \dots, n$$

and only one time along the **forward evaluation trace** for the m -dimension output values

$$y_1, y_2, \dots, y_m$$

- Get the partial derivatives of **all outputs** respect to **one variable** after computing along the derivative trace once
- **High efficiency** when $n \ll m$
- **Low efficiency** when $n \gg m$, which is common in machine learning and deep learning

2.3 Automatic differentiation - Reverse Mode

- Motivated by the **low efficiency** of *forward mode* when the derivatives of multiple variables are interested
 - In machine learning and deep learning, usually the dimension of input parameters are huge (such as 10^6) while the dimension of output is usually 1 (focusing on the loss function which produce a scalar)
- To find the partial derivative respect to i -th variable $\frac{\partial f}{\partial x_i}$ in the function

$$f = f(x_1, x_2, \dots, x_n)$$

we can always decompose the function into simple operation nodes

$$v_1, v_2, \dots, v_p$$

such that we can find some nodes $v_c, c \in [1, M]$ to compose the derivative

$$\frac{\partial y}{\partial x_i} = \frac{\partial y}{\partial v_{c_q}} \frac{\partial v_{c_q}}{\partial v_{c_{q-1}}} \dots \frac{\partial v_{c_2}}{\partial v_{c_1}} \frac{\partial v_{c_1}}{\partial x_i}$$

Denote

$$\bar{v}_c = \frac{\partial y}{\partial v_c}$$

The idea is to first compute following the **forward evaluation trace** same as that in *forward mode* while recording the **relationship** (such as the differential equation) between nodes, then sequentially calculate

$$\bar{v}_{c_q} = \frac{\partial y}{\partial v_{c_q}}, \bar{v}_{c_{q-1}} = \bar{v}_{c_q} \frac{\partial v_{c_q}}{\partial v_{c_{q-1}}}, \dots, \frac{\partial f}{\partial x_i} = \bar{v}_{c_1} \frac{\partial v_{c_1}}{\partial x_i}$$

- The example is still to find the partial derivatives at $x_1 = 2, x_2 = 5$ of the function

$$f(x_1, x_2) = \ln(x_1) + x_1 x_2 - \sin(x_2)$$


Similarly, decompose the function into simple operation nodes

$$\begin{aligned} v_{-1} &= x_1 \\ v_0 &= x_2 \\ v_1 &= \ln v_{-1} \\ v_2 &= v_{-1} \times v_0 \\ v_3 &= \sin v_0 \\ v_4 &= v_1 + v_2 \\ v_5 &= v_4 - v_3 \\ y &= v_5 \end{aligned}$$

The **forward evaluation computation graph** would be the same of that in *forward mode*. Here show the **reverse adjoint computation graph**

```
[36]: f = Digraph('ComputationGraph')
f.attr(rankdir='RL')
f.attr('node', shape='circle')
f.node('v_-1', label='<v<sub>-1</sub>>')
f.node('v_0', label='<v<sub>0</sub>>')
f.node('v_1', label='<v<sub>1</sub>>')
f.node('v_2', label='<v<sub>2</sub>>')
f.node('v_3', label='<v<sub>3</sub>>')
f.node('v_4', label='<v<sub>4</sub>>')
f.node('v_5', label='<v<sub>5</sub>>')
f.edge('v_1', 'v_-1')
f.edge('v_2', 'v_-1')
f.edge('v_2', 'v_0')
f.edge('v_3', 'v_0')
f.edge('v_4', 'v_1')
f.edge('v_4', 'v_2')
f.edge('v_5', 'v_4')
f.edge('v_5', 'v_3')
f
```

[36]:



- First, compute along the **forward evaluation trace**

$$\begin{array}{lll}
 v_{-1} = & x_1 = & 2 \\
 v_0 = & x_2 = & 5 \\
 v_1 = & \ln v_{-1} = & \ln 2 \\
 v_2 = & v_{-1} \times v_0 = & 2 \times 5 \\
 v_3 = & \sin v_0 = & \sin 5 \\
 v_4 = & v_1 + v_2 = & 0.693 + 10 \\
 v_5 = & v_4 - v_3 = & 10.693 + 0.959 \\
 y = & v_5 = & 11.652
 \end{array}$$

Then compute along the **reverse adjoint trace**

$$\begin{aligned}
 \bar{v}_5 &= \frac{\partial y}{\partial v_5} = 1 \\
 \bar{v}_4 &= \bar{v}_5 \frac{\partial v_5}{\partial v_4} = \bar{v}_5 \times 1 = 1 \\
 \bar{v}_3 &= \bar{v}_5 \frac{\partial v_5}{\partial v_3} = \bar{v}_5 \times (-1) = -1 \\
 \bar{v}_1 &= \bar{v}_4 \frac{\partial v_4}{\partial v_1} = \bar{v}_4 \times 1 = 1 \\
 \bar{v}_2 &= \bar{v}_4 \frac{\partial v_4}{\partial v_2} = \bar{v}_4 \times 1 = 1 \\
 \bar{v}_0 &= \bar{v}_2 \frac{\partial v_2}{\partial v_0} + \bar{v}_3 \frac{\partial v_3}{\partial v_0} = \bar{v}_2 \times v_{-1} + \bar{v}_3 \times \cos v_0 = 1.716 \\
 \bar{v}_{-1} &= \bar{v}_1 \frac{\partial v_1}{\partial v_{-1}} + \bar{v}_2 \frac{\partial v_2}{\partial v_{-1}} = \bar{v}_1 \times \frac{1}{v_{-1}} + \bar{v}_2 \times v_0 = 5.5
 \end{aligned}$$

Therefore, the result is

$$\begin{aligned}
 \frac{\partial f}{\partial x_1} &= \bar{v}_{-1} = 5.5 \\
 \frac{\partial f}{\partial x_2} &= \bar{v}_0 = 1.716
 \end{aligned}$$

- **Comments:**

- For a function $f : \mathbb{R}^n \Rightarrow \mathbb{R}^m$ with input $\mathbf{x} \in \mathbb{R}^n$ and output $\mathbf{y} \in \mathbb{R}^m$, needs to calculate m times along the **reverse adjoint trace**
- Get the partial derivatives of **one output** respect to **all variables** after computing along the reverse adjoint trace once
 - * Fits the idea of calculating **gradients**
- **High efficiency** when $n \gg m$
- **Low efficiency** when $n \ll m$
- Frequently used in ML/DL tasks and is supported by PyTorch, Tensorflow, ...

2.4 Computation graph and AutoGrad in PyTorch

- As mentioned above, PyTorch supports **reverse mode automatic differentiation**
 - PyTorch automatically creates a computation graph if **requires_grad=True**
 - For a given variable with **requires_grad=True**, identify the operations and record the required information for reverse adjoint trace computation
 - Can be checked by viewing the **requires_grad** attributes of independent variables
 - Use the **make_dots** method from **torchviz** module to visualize the computation graph
 - * Need support from **graphviz**
- Some comments:
 - PyTorch's computational graphs are **dynamic**, in the sense that a new graph is created in each forward pass. On the other hand, the computational graphs constructed by Tensorflow are **static**, in the sense that the same graph is used over and over in all iterations during training
 - In general, static graphs are more efficient because it only needs to be optimized once. Optimization generally consists of distributing the computations over the graph nodes

- across multiple GPUs if more than one GPU is available, or just fusing some nodes of the graph if the result logic won't be impacted by such fusion
- However, static graphs do not lend themselves well to recurrent neural computations because the graph itself can change from iteration to iteration

```
[37]: # The default setting is requires_grad = False
```

```
x = torch.tensor(5.)  
y = 3 * x**2 + x  
print(x, x.requires_grad)  
print(y, y.requires_grad)
```

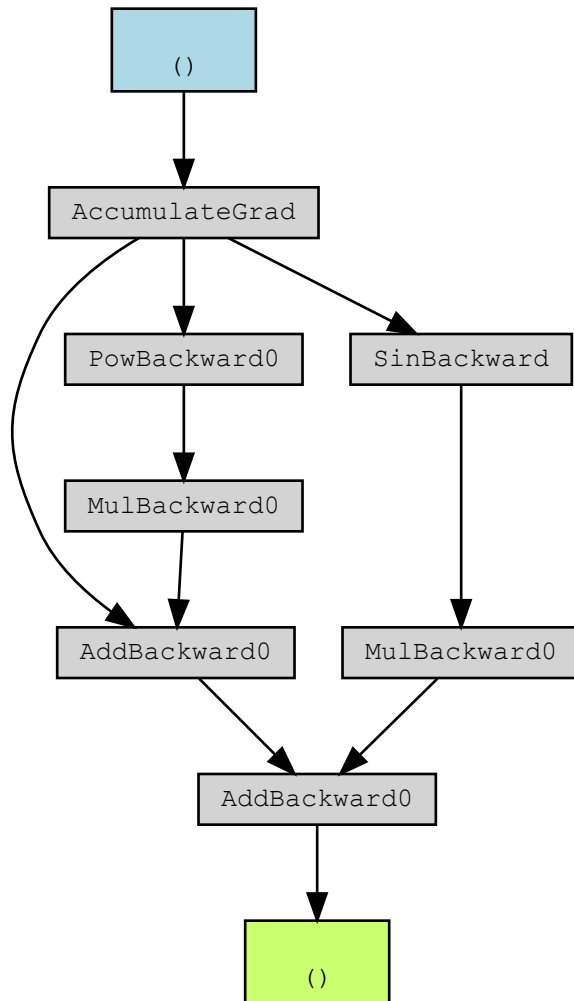
```
tensor(5.) False  
tensor(80.) False
```

```
[38]: # Use the torchviz module for visualization
```

```
from torchviz import make_dot  
  
x = torch.tensor(5., requires_grad=True)  
y = 3 * x**2 + x + 4 * torch.sin(x)  
# Independent variable would show the requires_grad attribute  
print(x, x.requires_grad)  
# Dependent variable would show the grad_fn attribute  
print(y, y.requires_grad)  
make_dot(y)
```

```
tensor(5., requires_grad=True) True  
tensor(76.1643, grad_fn=<AddBackward0>) True
```


```
[38]:
```



- Notice the **grad_fn** attribute. Tensor use this attribute to do the backwards computation.
- When the dimension of input tensor is greater than one, the dimension of each step would also show on the graph
 - Very useful when building neural networks

```
[39]: x = torch.randn((3, 4, 4), requires_grad=True)
      y = x**2 + torch.sin(x)
      make_dot(y)
```

[39]:



- One can access the gradient by calling **backward()** method
- After invoke **backward()** method on the dependent variable we are interested in, the partial derivative would automatically be returned to the **grad** attribute of variables contributed to that dependent variable
- For **independent** variables, we can directly call **backward()** method and see the gradients
- In the following case:

$$z = x_1^2 + x_2^2$$

The partial derivatives at $x_1 = 5, x_2 = 2$ are

$$\frac{\partial z}{\partial x_1} = 2x_1 = 10, \frac{\partial z}{\partial x_2} = 2x_2 = 4$$

```
[40]: x_1 = torch.tensor(5., requires_grad=True)
x_2 = torch.tensor(2., requires_grad=True)
z = x_1**2 + x_2**2
z.backward()
print(x_1, x_1.grad)
print(x_2, x_2.grad)
print(z)
```

```
tensor(5., requires_grad=True) tensor(10.)
tensor(2., requires_grad=True) tensor(4.)
tensor(29., grad_fn=<AddBackward0>)
```

- For **intermediate dependent** variables (*non-leaf tensor* in pytorch documentation) used in computation, the gradient value is usually not considered and would be cleared after the

computation along reverse adjoint trace.

- If the value is indeed needed, the `retain_grad()` method should be invoked

```
[41]: x_1 = torch.tensor(5., requires_grad=True)
x_2 = torch.tensor(2., requires_grad=True)
y = 0.5 * x_2**2
y.retain_grad()
z = x_1**2 + y
z.backward()
print(x_1, x_1.grad)
print(x_2, x_2.grad)
print(y, y.grad)
print(z)
```

```
tensor(5., requires_grad=True) tensor(10.)
tensor(2., requires_grad=True) tensor(2.)
tensor(2., grad_fn=<MulBackward0>) tensor(1.)
tensor(27., grad_fn=<AddBackward0>)
```

- It needs to be noticed that the gradients accumulate when calling `backward()`

```
[42]: x = torch.tensor(5., requires_grad=True)
for ii in range(2):
    y = 3 * x**2
    y.backward()
    print(x, x.grad)
    print(y)
```

```
tensor(5., requires_grad=True) tensor(30.)
tensor(75., grad_fn=<MulBackward0>)
tensor(5., requires_grad=True) tensor(60.)
tensor(75., grad_fn=<MulBackward0>)
```

- Therefore, when calculating the gradients repeated in loops, we usually need to **zero** the gradients before calling `backward()` method, by calling `zero_()` method

```
[43]: x = torch.tensor(5., requires_grad=True)
for ii in range(2):
    try:
        x.grad.zero_()
    except Exception as e:
        print(e)
    y = 3 * x**2
    y.backward()
    print(x, x.grad)
    print(y)
```

```
'NoneType' object has no attribute 'zero_'
tensor(5., requires_grad=True) tensor(30.)
tensor(75., grad_fn=<MulBackward0>)
```

```
tensor(5., requires_grad=True) tensor(30.)
tensor(75., grad_fn=<MulBackward0>)
```

- Sometimes, we would need to call the **backward()** method multiple times, like using MSE loss and Cross-Entropy loss separately for localization and detection in CV. In this case, **retain_graph** should be set to **True** except in the last **backward()** method if any **intermediate dependent variables** are used.

```
[44]: x = torch.tensor(5., requires_grad=True)
      y = x**2
      z_1 = x + y
      z_2 = x**2 + y
      z_3 = x**3 + y
      z_1.backward(retain_graph=True)
      z_2.backward(retain_graph=True)
      z_3.backward()
      print(x, x.grad)
```

```
tensor(5., requires_grad=True) tensor(116.)
```

- Generally speaking, PyTorch can compute gradients for any number of parameters and any complex functions

```
[45]: x = torch.arange(5.).requires_grad_(True)
      y = torch.sum(x**2)
      y.backward()
      print(x)
      print(y)
      print(x.grad)
```

```
tensor([0., 1., 2., 3., 4.], requires_grad=True)
tensor(30., grad_fn=<SumBackward0>)
tensor([0., 2., 4., 6., 8.])
```

```
[46]: x = torch.arange(5.).requires_grad_(True)
      y = torch.mean(torch.log(x**2 + 1) + 5 * x)
      y.backward()
      print(x)
      print(y)
      print(x.grad)
```

```
tensor([0., 1., 2., 3., 4.], requires_grad=True)
tensor(11.4877, grad_fn=<MeanBackward0>)
tensor([1.0000, 1.2000, 1.1600, 1.1200, 1.0941])
```

3 Simple Gradient Descent example

- To show how automatic gradient in PyTorch and the simple gradient descent algorithm work, here are two examples where the input and output are both 1-dimension variables
- For convenience, define the gradient descent algorithm as a function and the visualization process as a function

```
[47]: import matplotlib.pyplot as plt
import numpy as np

# Function of simple gradient descent
def gradient_descent(objective, step_size=0.05, max_iter=100, init=0):
    # Initialize
    x_hat = torch.tensor(init, dtype=torch.float32, requires_grad=True)
    # Record the value iteration process
    x_hat_arr = [x_hat.detach().numpy().copy()]
    obj_arr = [objective(x_hat).detach().numpy()]
    # Iterate
    for ii in range(max_iter):
        # Compute gradient
        if x_hat.grad is not None:
            x_hat.grad.zero_()
        out = objective(x_hat)
        out.backward()
        # Update x_hat
        # Stop tracking gradients
        with torch.no_grad():
            x_hat -= step_size * x_hat.grad
        x_hat_arr.append(x_hat.detach().numpy().copy())
        obj_arr.append(objective(x_hat).detach().numpy())
    return np.array(x_hat_arr), np.array(obj_arr)

# Function to visualize iteration process
def visualize_result(x_arr, obj_arr, objective, x_true=None, vis_arr=None):
    # The horizontal range of figure
    if vis_arr is None:
        vis_arr = np.linspace(np.min(x_arr), np.max(x_arr))
    fig, ax = plt.subplots(figsize=(8, 3), dpi=100)
    ax.plot(vis_arr, [objective(torch.tensor(x)).numpy() for x in vis_arr],
            label='Objective')
    ax.plot(x_arr, obj_arr, 'o-', label='Gradient Steps')
    # If true minimum location is provided, plot it out
    if x_true is not None:
        ax.plot(np.ones(2) * x_true, plt.ylim(), label='True x')
    ax.plot(np.ones(2) * x_arr[-1], plt.ylim(), label='Final x')
    plt.legend(loc='lower right')
```

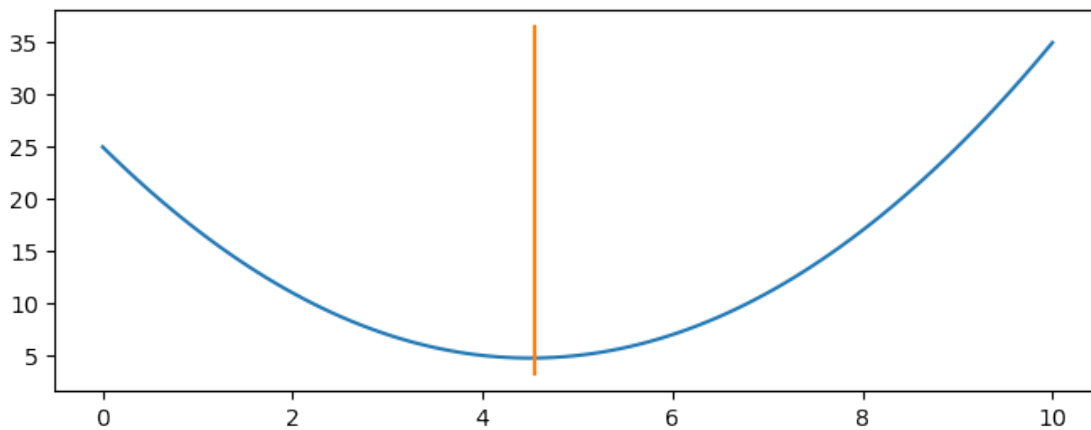
```
plt.show()
```

3.1 Convex function - Converge to global minimum

- Here is an example where the objective function is very simple and gradient descent can easily reach the global minimum
- First plot out the objective function

```
[48]: # Objective function
def Objective(x):
    return torch.abs(x) + (x - 5)**2

x = torch.linspace(0, 10, 100)
y = Objective(x)
x_true = float(x[np.argmin(y)])
fig, ax = plt.subplots(figsize=(8, 3), dpi=100)
# Need to convert the tensor to numpy arrays for plots
ax.plot(x.numpy(), y.numpy())
# Show the minimum location
ax.plot(x_true * np.ones(2), plt.ylim())
plt.show()
```

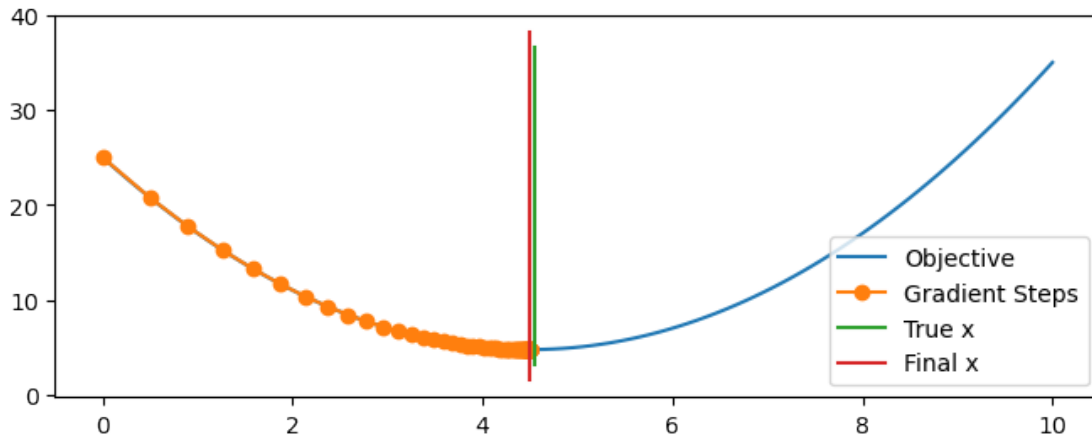


- Apply simple gradient descent on this function

```
[49]: # Call the functions to get the minimum
x_hat_arr, obj_arr = gradient_descent(Objective,
                                     step_size=0.05,
                                     max_iter=150,
                                     init=0)

visualize_result(x_hat_arr,
                obj_arr,
```

```
Objective,
x_true=x_true,
vis_arr=np.linspace(0, 10, num=100))
```



- It seems that the final value approximately reaches the global minimum. Let's check the result by printing them out

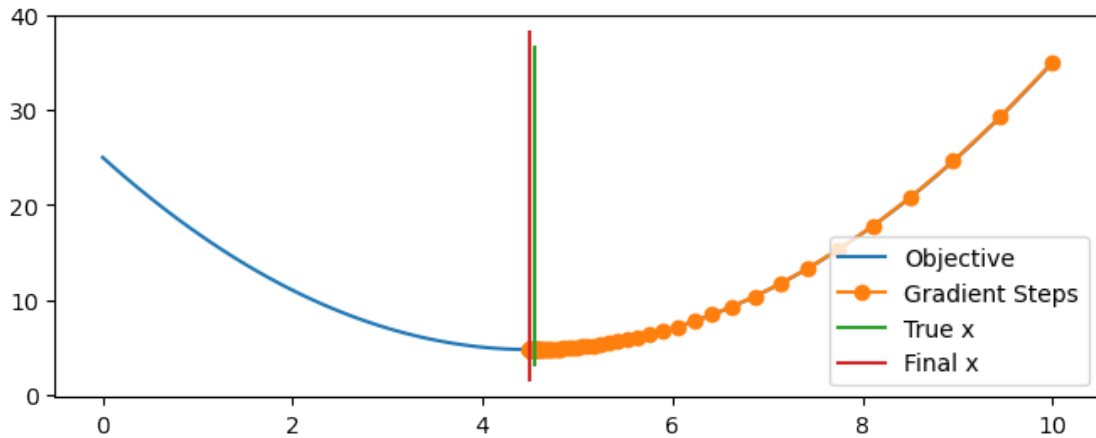
```
[50]: print("The true minimum location is:          %f" % (x_true))
      print("The final location got by GD is:        %f" % (x_hat_arr[-1]))
      print("Errors between true and final location is: %f" %
            (x_true - x_hat_arr[-1]))
```

```
The true minimum location is:          4.545455
The final location got by GD is:        4.499998
Errors between true and final location is: 0.045456
```

- In this case, the start point does **not** matter

```
[51]: # Call the functions to get the minimum
x_hat_arr, obj_arr = gradient_descent(Objective,
                                     step_size=0.05,
                                     max_iter=150,
                                     init=10)

visualize_result(x_hat_arr,
                 obj_arr,
                 Objective,
                 x_true=x_true,
                 vis_arr=np.linspace(0, 10, num=100))
```

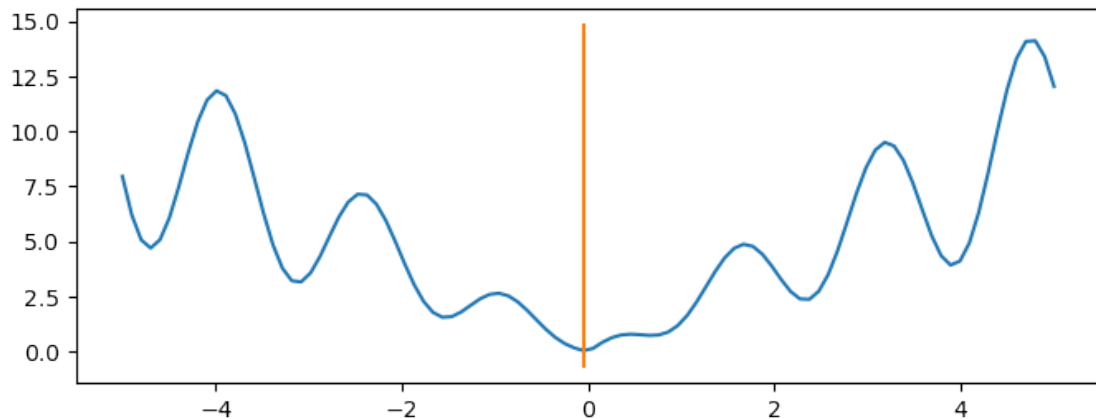


3.2 Non convex function - Converge to local minimum

- Here is an example where the objective function is a bit complex and gradient descent falls into local minimum
- First plot out the objective function

```
[52]: def Objective(x):
        return x * torch.cos(4 * x) + 2 * torch.abs(x)

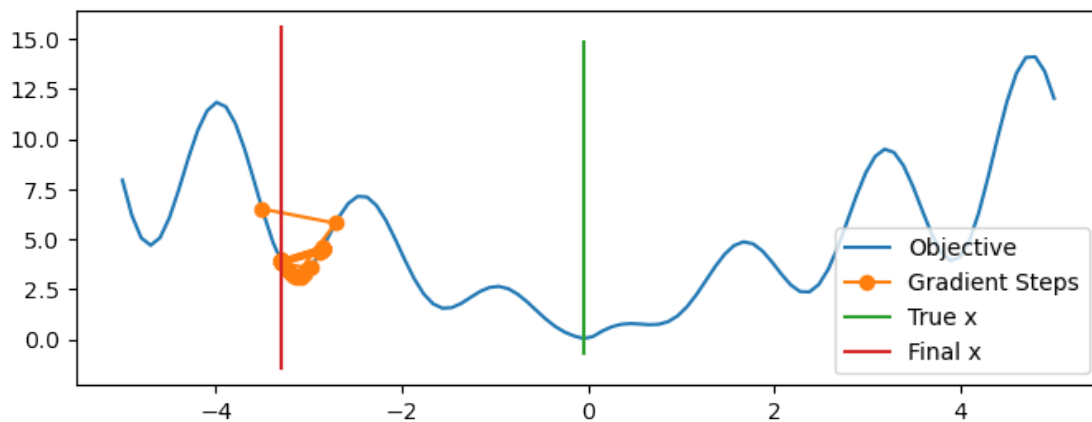
x = torch.linspace(-5, 5, 100)
y = Objective(x)
x_true = float(x[np.argmin(y)])
fig, ax = plt.subplots(figsize=(8, 3), dpi=100)
ax.plot(x.numpy(), y.numpy())
ax.plot(x_true * np.ones(2), plt.ylim())
plt.show()
```



- As what was done before, apply simple gradient descent to it

```
[53]: x_hat_arr, obj_arr = gradient_descent(Objective,
                                           step_size=0.05,
                                           max_iter=150,
                                           init=-3.5)

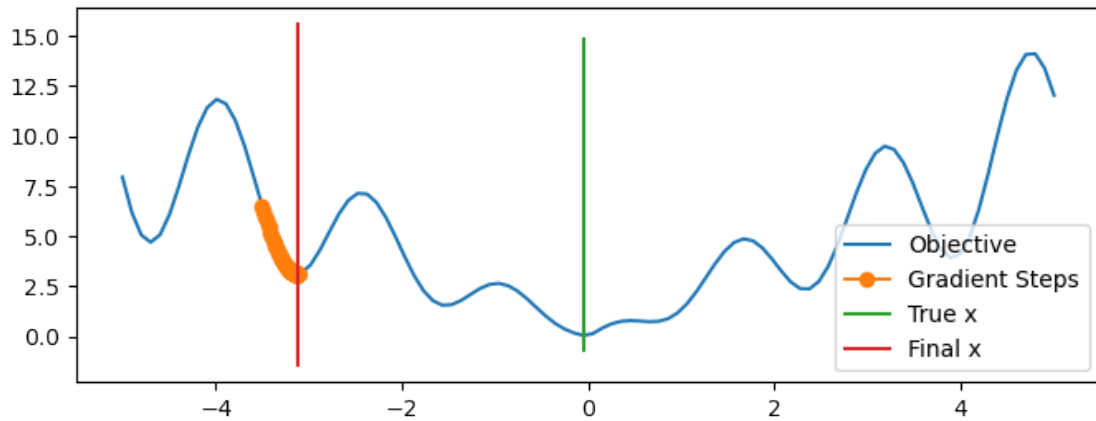
visualize_result(x_hat_arr,
                 obj_arr,
                 Objective,
                 x_true=x_true,
                 vis_arr=np.linspace(-5, 5, num=100))
```



- In this case, it seems that the **step size** is so huge that the oscillation occurs and gradient descent even cannot reach a local minimum. Try with a smaller step size

```
[54]: x_hat_arr, obj_arr = gradient_descent(Objective,
                                           step_size=0.001,
                                           max_iter=150,
                                           init=-3.5)

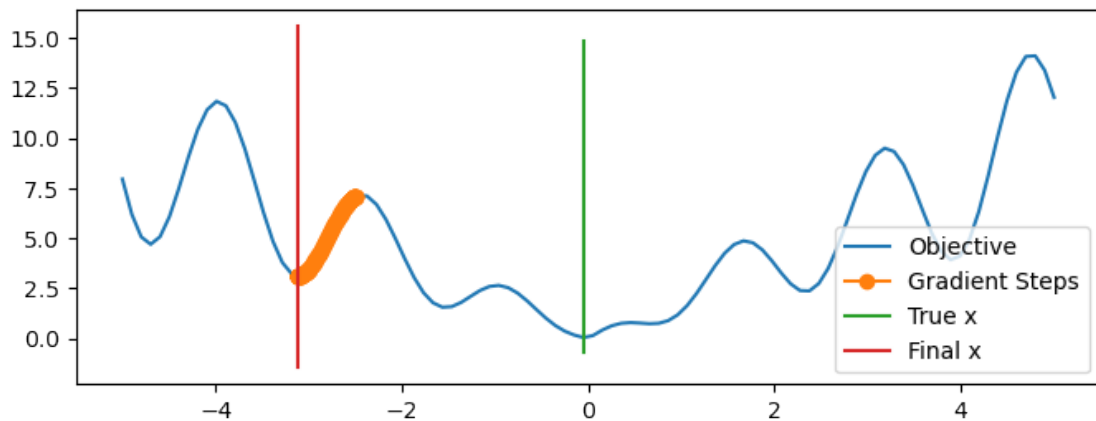
visualize_result(x_hat_arr,
                 obj_arr,
                 Objective,
                 x_true=x_true,
                 vis_arr=np.linspace(-5, 5, num=100))
```



- With a smaller step size, it successfully converge to a local minimum. However, we would always look forward to a global minimum.
- In this case, the **start point will affect final convergence**. Try with different start points

```
[55]: x_hat_arr, obj_arr = gradient_descent(Objective,
                                           step_size=0.001,
                                           max_iter=150,
                                           init=-2.5)

visualize_result(x_hat_arr,
                 obj_arr,
                 Objective,
                 x_true=x_true,
                 vis_arr=np.linspace(-5, 5, num=100))
```



```
[56]: x_hat_arr, obj_arr = gradient_descent(Objective,
                                           step_size=0.001,
```

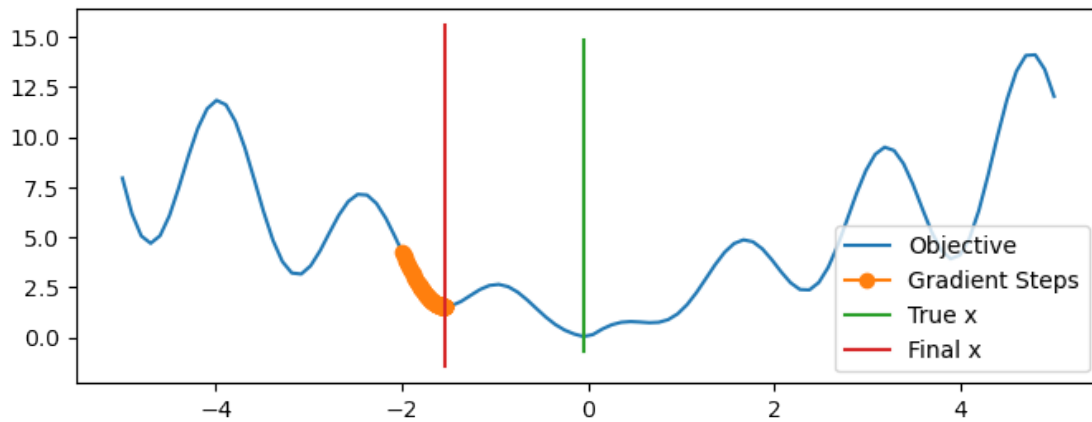


```

max_iter=150,
init=-2)

visualize_result(x_hat_arr,
                 obj_arr,
                 Objective,
                 x_true=x_true,
                 vis_arr=np.linspace(-5, 5, num=100))

```

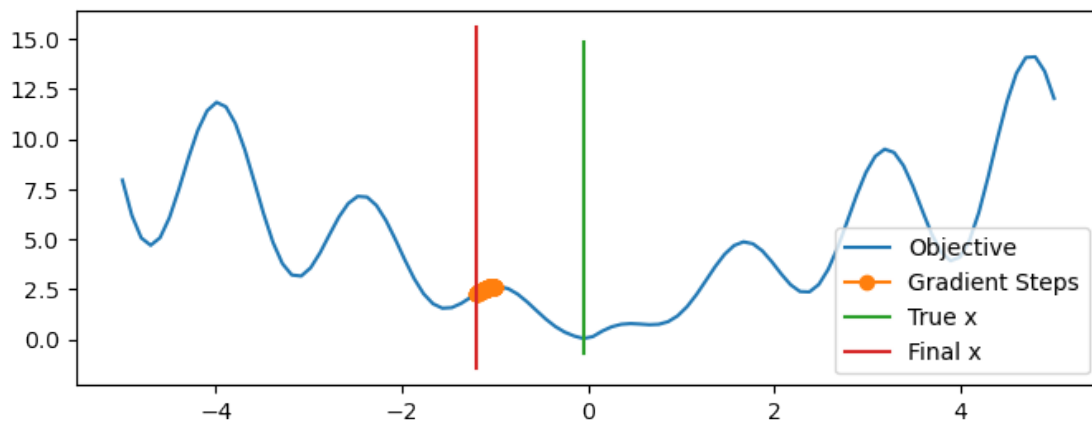


```

[57]: x_hat_arr, obj_arr = gradient_descent(Objective,
step_size=0.001,
max_iter=150,
init=-1)

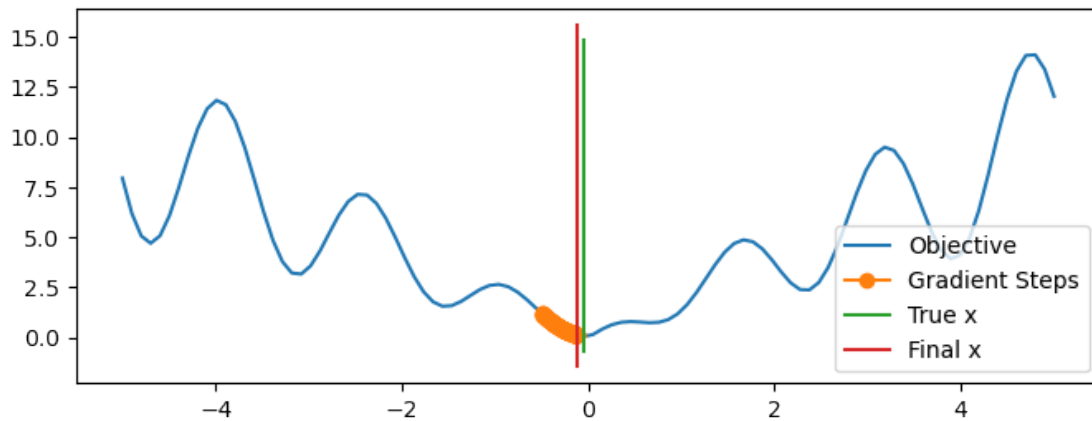
visualize_result(x_hat_arr,
                 obj_arr,
                 Objective,
                 x_true=x_true,
                 vis_arr=np.linspace(-5, 5, num=100))

```



```
[58]: x_hat_arr, obj_arr = gradient_descent(Objective,
                                           step_size=0.001,
                                           max_iter=150,
                                           init=-0.5)

visualize_result(x_hat_arr,
                 obj_arr,
                 Objective,
                 x_true=x_true,
                 vis_arr=np.linspace(-5, 5, num=100))
```



- After several trials, we finally reaches the approximated location of global minimum. However, in complex optimization problems, we will not be able to select a best start point so easily by visualization, which **makes finding the global minimum very difficult**.